

Jev, Vibe Check 🎧

Decision Offloading


What does an agent gain
when judgment is offloaded?



Jev
by TypeSafe AI



GEODE
Self-evolve agent harness



Decision offloading: design and validation


What does an agent gain by offloading its verdicts? · Korean 13:29, English 13:32

01	Question / system	13:43
	Jev public surface · Research question · GEODE and judge replacement	
02	Implementation / pilot	14:54
	Development scaffold · Pilot observations · Public surface and research	
03	Design / methods	17:02
	Comparison and completion criteria · Data and generation · Procedure and tracks · Policy freeze and analysis · Metric reading guide	
04	Quality by task	20:30
	Completion / Choice · Joint judgment / Noul · Intent and target · Candidate selection / Score · External trajectory judgment · External candidate selection	
05	Time / cost	21:38
	Main study / verifier latency · Calls and accounting scope	
06	Signals / policy	21:59
	Repeat consistency · Error ranking and calibration · Frozen conditional offloading	
07	Execution	22:44
	Repeated runs · fault injection · repair and delivery	
08	Use / reproduction	25:57
	Operating conditions · Public evidence · Reader review paths · Sources and production	
09	Conclusion	26:48
	Findings and adoption	

Decision offloading: design and validation

What does an agent gain by offloading its verdicts? · Korean 13:29, English 13:32

01	Question / system	13:43
	Jev public surface · Research question · GEODE and judge replacement	
02	Implementation / pilot	14:54
	Development scaffold · Pilot observations · Public surface and research	
03	Design / methods	17:02
	Comparison and completion criteria · Data and generation · Procedure and tracks · Policy freeze and analysis · Metric reading guide	
04	Quality by task	20:30
	Completion / Choice · Joint judgment / Noul · Intent and target · Candidate selection / Score · External trajectory judgment · External candidate selection	
05	Time / cost	21:38
	Main study / verifier latency · Calls and accounting scope	
06	Signals / policy	21:59
	Repeat consistency · Error ranking and calibration · Frozen conditional offloading	
07	Execution	22:44
	Repeated runs · fault injection · repair and delivery	
08	Use / reproduction	25:57
	Operating conditions · Public evidence · Reader review paths · Sources and production	
09	Conclusion	26:48
	Findings and adoption	



SOUNDTRACK
Nemzzz
LOCKED IN · drill · 136–151 BPM

01	GASS feat. Travis Scott	140 BPM
	■□□□ ■□□□ ■□■□ ■□■□	0:00–1:54
02	KILLIN EM	144 BPM
	■□■□ ■□■□ ■□■□ ■□■□	1:54–3:47
03	TODAY	142 BPM
	□□□□ □□■□ □□□□ ■□□□	3:47–5:50
04	TOKYO DRIFT	136 BPM
	■□■□ □□■□ □□■□ ■□□□	5:50–7:38
05	PRINCE OF THE SCENE	149 BPM
	■□□□ ■□□□ ■□□□ □□□□	7:38–9:34
06	4AM	151 BPM
	■□□□ ■□□□ ■□□□ ■□□□	9:34–12:10

First-cycle times · measured cues, build-specific reading time
After the six tracks, repeat from GASS

Three primitives for decisions in code



Jev · typed judgments and probabilities

State and typed questions in; typed answers and probabilities out. Code decides the next action.

QUESTION TYPE

- Choice · select an option
- Noul · yes or no
- Score · ordered levels

TYPED ANSWER

choice + probabilities
confidence

label: P(.)

...

p

$P(\text{yes}) \in [0, 1]$

Σ

$\text{score} = \Sigma i \times P(i)$

legend · probabilities · confidence

CODE CONSUMER

- Choose a branch
- ?

Check a condition
- Rank by rubric

Jev’s public surface: an API for decisions

POST /v1/systemone

Shared context and independent questions in one request

model	Select the model	
state	Evidence · request · candidate, once	
questions	verdict	Choice
	evidence	Noul
	quality	Score



Answers keyed by question

Type-specific shapes, with no illustrative values

- Choice**
Selected option + probabilities + confidence
- Noul**
P(yes) ∈ [0, 1]
No separate confidence
- Score**
 $\Sigma \text{ index} \times P(\text{index}) \cdot \text{probabilities} \cdot \text{confidence}$

What the caller still owns Acceptance criteria · next action · retry and repair budgets

Batch independent questions; put answer-dependent questions in a later call.
API shapes, not next-token probabilities or end-to-end task success rates.

Jev's public surface: an API for decisions

API · SDK

Published contract

state · model · questions → typed answers

Request shapes, answer spaces and probability fields are documented.

Launch post · founder

Vendor account

New architecture · parallel sampler · RLCD

Novelty is a vendor claim; no reproducible implementation was found.

Network · training recipe

Not disclosed in reviewed sources

Backbone · attention mask · loss · weights

Reuse of an encoder or decoder from a specific LLM remains unverified.

When can a faster judgment be delegated?

Final native evidence · editorial interpretation

Original runs preserved; corrections labeled separately

Offloading depended on the judgment and delivery contract.

Fixed intent classification

Does it identify the correct intent and target?
Separate panel items from actual tasks

Noul joint · candidate selection

Does quality hold for joint conditions and candidate ranking?
Each primitive has its own comparison unit

Probability · conditional offloading


Can uncertain judgments go to a stronger judge?
Freeze policy; evaluate panel and execution separately

Next: matched-input quality, signals that rank errors, then actual delivery.

Who builds it: and what proves it works?


From a human question to verifiable evidence

Coding hosts ≠ experimental arms




GEODE’s execution map: rounds and final verification

Request · session · execution scope · remaining budget
Prepare input once; repeat rounds within one physical turn.




GEODE's execution map: rounds and final verification



Where Jev fits: judgments remain separate from actions

Keep the main execution loop

Current core Jev routes require explicit selection



Choice · completion verdict

Input	Root request · candidate · retained evidence
GEODE map	passed · should_retry · feedback
Reader	Finalizer · deliver / repair / escalate

GEODE retains execution authority and repair budgets

Choice · evidence classification

Input	Observations + bounded cognitive state
Write	Append a code-owned next check to subgoals
Reader	Next Reflection · state snapshots

Not belief/confidence synthesis or direct root-prompt injection

Reflection after the tool checkpoint

Evaluation-only · opt-in

Choice intent + target classification

analyze_request returns to root context; no action execution or permission

Where Jev fits: judgments remain separate from actions

Keep the main execution loop

Current core Jev routes require explicit selection



Comparisons at design time, separate from main paths

Score

Selection quality on the same four frozen candidates

Current candidate_sampling selection uses an LLM

Noul

Judge contradictions and missing evidence separately

Measure joint judgments and subsequent repair

Cascade

Escalate only uncertain Jev judgments to Astra

Compare thresholds, call cost and false acceptance

Keep global LLM judgment; swap only assigned slots


Evaluation-only · opt-in

Choice intent + target classification

analyze_request returns to root context; no action execution or permission

One contract, two coding-host contexts

HOST · System instructions / tools / permissions and sandbox



AGENTS.md: a map to the right instructions

From shared rules to relevant evidence

AGENTS.md

Scope · ownership · acceptance · verification
Codex reads it; CLAUDE.md imports @AGENTS.md

- docs/workflow.md → task owner
Route the request to its code and document owner
- .agents/skills/ → SKILL.md
.claude/skills/ aliases the same source files
- Selected references → code, callers and tests
Read the chosen skill fully, then follow the evidence

Routes selected for this project

Design and measurement

[geode-eval](#) · [stanford-test-time-compute](#)
05 preregistration + 06 reliability → freeze / admission

Build and integration

[geode-workflow](#) · [geode-gitflow](#)
Bounded role contracts → patches → owner verification

Film and review

[technical-deck-report](#) · [viz-frame-audit](#)
DESIGN.md + fluent-korean → KO/EN render review


- Delegation

Role, inputs, edit scope, acceptance → spawn → patch / findings
- Iteration

Owner reproduces and checks → findings back → revise and recheck

GEODE.md + .geode/skills/ supply runtime context; the host enforces permissions and sandboxing.

Worktree lifecycle, from isolation to merge and cleanup



Uncertain merge outcome → verify remote state before retry or cleanup

ADMISSION Freeze a verified SHA, inputs and policy; pass admission to run. Main promotion follows separate approval.

A faster verdict, but why did one task take longer?

○ LLM ■ Jev One line = one matched task–repetition pair

Final-judge call · seconds

explicit · r0


explicit · r1

context · r0 *

context · r1

korean · r0

korean · r1




Judge median: 7.235 s → 0.426 s

A faster verdict, but why did one task take longer?

○ LLM ■ Jev One line = one matched task–repetition pair

Final-judge call · seconds


GEODE runtime · s (setup and oracle excluded)



Judge median: 7.235 s → 0.426 s

Runtime median: 50.198 → 46.439 s · both 6/6 pass

A faster verdict, but why did one task take longer?




verdict -6.07 · extra reflection +7.88 · first reflection -2.29 · main +1.21 · rest +0.01 → runtime +0.74 s

One more Cognitive reflection ran right before the verdict: its 7.88 s outweighed the 6.07 s the verdict saved.

Cadence: Astra Cognitive confidence ≥ 0.8 doubles the interval. Logged: LLM 0.8; Jev 0.65 after the extra pass. First-pass confidence was not logged; the cause of the extra pass remains unresolved.

Six verdicts: 18.2× speedup, what about total time?

Six-pair sums · Jev – LLM · seconds · Natural runs r2 natural (3 inboxes × 2 repeats)



One verdict · explicit/r0

7.661 $\rightarrow$ 0.419 s

×18.3

Astra API-equiv. \$0.0398 · Jev list price \$0.00011

Six verdicts

45.779 → 2.515 s

×18.2

\$0.2448 · \$0.00065 · 14.95% of baseline time

Whole task · six pairs

306.244 → 276.465 s

×1.11

Astra-equiv. \$1.077 $\rightarrow$ \$0.900 .

Jev separate · billing unknown

Holding other stages fixed: a zero-time verdict caps savings at $14.95\% \cdot \text{ratio} = \text{LLM} \div \text{Jev}$

Verdicts -43.264 s + other stages +13.485 s = total -29.779 s (-9.72%).

Call order: in all twelve runs its calls finished before the verdict call, and no request went back after the verdict.

The increase was observed in Astra execution before the Jev call:

10.5 s from korean/r1 (first call, same 3,011 input tokens: 9.6 → 14.3 s; both calls' output 512 → 723 tokens)

The same answer, different observations?

The request and candidate “Order A-104 has shipped” stay fixed across three tool observations.

Supplied observation	LLM Jev · actual verdict	Code-projected branch
status: shipped Supports the same candidate	supported supported Jev confidence 1.00	Pass → no revision
status: cancelled Contradicts the same candidate	contradicted contradicted Jev confidence 1.00	Not passed → request revision
observations: [] No evidence for the candidate	insufficient_evidence insufficient_evidence Jev confidence 0.95	Not passed → request revision

SCOPE

Both engines matched 3/3 labels. Repair branches are code projections; this diagnostic made no downstream agent call.

Missing evidence, a “contradiction” verdict

Same input · A-104 candidate: shipped · Lookup observations: 0

LLM Astra / xhigh

Request contract · one field: verdict

Actual return

```
{"verdict": "insufficient_evidence"}
```

Response latency 6.607 s · 929 input / 106 output tokens

...step-2-025475da

Jev Jev 1.13

Request contract · Choice primitive

Actual return

```
{"verdict": {  
  "type": "choice", "choice": "contradicted",  
  "confidence": 0.58,  
  "probabilities": {"supported": 0.0,  
    "contradicted": 0.72,  
    "insufficient_evidence": 0.28}  
}}
```


Response latency 8.506 s · 933 input / 51 output tokens

...step-2-a1a1d069

CONSEQUENCE Jev returned contradicted instead of insufficient_evidence. Both labels routed the run to repair.

The cost from a failed verdict to recovery

Before lookup: the answer was shipped, with no recorded evidence.



Answer: shipped → shipped · The new lookup supports the same answer.

Recovery after the first negative verdict	LLM	Jev
Replan + main agent + Cognitive + rejudge	1 + 2 + 2 + 1	1 + 2 + 2 + 1
Until the final verdict completes	58.855 s	50.321 s
Recorded input / output tokens	8,003 / 1,405	8,030 / 1,363
Astra API equivalent / Jev tariff estimate	\$0.150280 / no Jev	\$0.135560 / \$0.000044058

Had we read the confidence, what would it have flagged?

All 13 recorded Jev verdicts, ordered by confidence. Execution used labels; confidence was logged only.

- agrees with the independent check
- miss

Natural task runs

task verifier · 6

Fixed-input judgments


expected label · 3

Fault injection first verdict

expected label · 2

Fault injection re-judgment

oracle · 2



POST HOC A cut below 0.60 would flag the one miss and two correct verdicts: 3 rechecks for 1 error.
Threshold selection still needs held-out evidence.

Check order A-104 status · Choice

LLM verdict | Astra | 1x | 588a1c8 | replay 00:00:00.000

```
[reconstructed] Harbor ATIF replay; not a raw PTY capture.
[task] Process every inbox item.
<inbox>[{"id": "status_one", "request": "What is the current status of A-104?",
"candidates": ["A-104"]}]/>
```

POST-RUN ORACLE: PASS | separate verifier receipt

Jev verdict | Astra | 1x | 588a1c8 | replay 00:00:00.000

```
[reconstructed] Harbor ATIF replay; not a raw PTY capture.
[task] Process every inbox item.
<inbox>[{"id": "status_one", "request": "What is the current status of A-104?",
"candidates": ["A-104"]}]/>
```

POST-RUN ORACLE: PASS | separate verifier receipt

Check order A-104 status · Choice

LLM verdict Astra 1x 588a1c8 replay 00:01:10.267	END	Jev verdict Astra 1x 588a1c8 replay 00:00:50.133	END / held
[reconstructed] Harbor ATIF replay; not a raw PTY capture. [task] Process every inbox item. <inbox>[{"id": "status_one", "request": "What is the current status of A-104?", "candidates": ["A-104"]}]/> \$ [lookup_order_status] {"items": [{"id": "status_one", "order_id": "A-104"}]} {"result": {"items": [{"id": "status_one", "order_id": "A-104", "status": "shipped"}]}} [agent] {"items": [{"id": "status_one", "order_id": "A-104", "status": "cancelled", "disposition": "answered"}]} [agent] {"items": [{"order_id": "A-104", "status": "shipped", "disposition": "answered", "id": "status_one"}]} CONTROLLED INPUT: {"items": [{"id": "status_one", "order_id": "A-104", "status": "cancelled", "disposition": "answered"}]} Judgment 2: supported 5311.9 ms CODE FEEDBACK: Confirm the final response retains those evidence boundaries. Feedback consumed: step-2-53687a5a (confirmed at call end) POST-RUN ORACLE: PASS separate verifier receipt	[reconstructed] Harbor ATIF replay; not a raw PTY capture. [task] Process every inbox item. <inbox>[{"id": "status_one", "request": "What is the current status of A-104?", "candidates": ["A-104"]}]/> \$ [lookup_order_status] {"items": [{"id": "status_one", "order_id": "A-104"}]} {"result": {"items": [{"id": "status_one", "order_id": "A-104", "status": "shipped"}]}} [agent] {"items": [{"id": "status_one", "order_id": "A-104", "status": "cancelled", "disposition": "answered"}]} [agent] {"items": [{"order_id": "A-104", "status": "shipped", "disposition": "answered", "id": "status_one"}]} CONTROLLED INPUT: {"items": [{"id": "status_one", "order_id": "A-104", "status": "cancelled", "disposition": "answered"}]} Judgment 2: supported 468.7 ms CODE FEEDBACK: Confirm the final response retains those evidence boundaries. Feedback consumed: step-2-2f43ab05 (confirmed at call end) POST-RUN ORACLE: PASS separate verifier receipt		

Replay ended · reading the results

How far can a verdict probability be trusted?

RLHF

Human preference between responses

Goal: Preferred responses

RLVR

Rewards verified by answers and tests

Goal: Passing verifiable checks

RLCD

Training signal not disclosed

Goal: Calibrated decision probabilities

RLCD's stated goal: align reported probabilities with observed frequencies.

How far can a verdict probability be trusted?

RLHF

Human preference between responses

Goal: Preferred responses

RLVR

Rewards verified by answers and tests

Goal: Passing verifiable checks


RLCD

Training signal not disclosed

Goal: Calibrated decision probabilities

RLCD's stated goal: align reported probabilities with observed frequencies.

External measurement · was a higher probability right more often? (JEV-as-a-Judge, CMU, 2026-09-22)



External-paper measurements; GEODE calibration and error ranking were evaluated separately.

CLM: contrastive learning for state–action pairs


PAIRED INPUTS

FROZEN BACKBONE


TRAINABLE HEADS

CONTRASTIVE · InfoNCE


state





Qwen3-8B
last-token pooling




state head



action



action head



state ↓

action →

+		
	+	
		+

Rows + columns: paired target
Illustration, not measured scores

Decoder backbone → input vectors · head training → state–action similarity

TRAINING

The public trainer targets recorded state–action pairs. Whether high similarity yields task success requires a separate evaluation.


Reusable action vectors for recurring candidates

PREPARE ONCE

Fixed candidates

Qwen → action head

Reuse vectors



scaled cosine

$\cos(s, a_i)$

$\vdots$

$\cos(s, a_n)$

→ softmax

→ Candidate probs.

$P(a_i | \text{pool})$

...


Within the supplied pool

EACH DECISION

New state + question

Qwen → state head

New vector



Compare each candidate with s

Reuse requires stable candidate text, preprocessing, backbone and heads.

COST SCOPE

The author's timing covers scoring with cached candidates. GEODE evaluation must also count candidate preparation, repeated scoring and wrong-branch repair.

Author's warm-cache result: 44 ms / 579 ms at 1,024 candidates · local CLM / hosted Jev

200-token state · 15 tokens/candidate · RTX 4090 · 5 trials · excludes initial candidate preparation

Change the judgment model, retain the execution owner

Comparison design / fixed-input evaluation and live execution

What changed, and what stayed fixed?

Separate a judgment on fixed input from the execution that consumes it

Changed

Judgment model

○ **A · GPT-6 Astra**

Codex subscription · effort xhigh

■ **B · Jev 1.13.0**

Direct API · typed responses

Fixed-input comparison

Matched states, pools or requests; the same task-specific criteria

Source, spec, input, policy and budget pinned before each run

E2E execution owner

Astra xhigh / GEODE in both arms

The root uses tools and repairs; intermediate observations and paths can differ.

Task-success evidence

Separate task verification plus execution-path and observation records

Judge acceptance $\neq$ task success; provider prompts are not byte-identical

The adopter defines completion; verifiers check it

Adopter design guide / this diagnostic contract

What observable evidence counts as completion?

Define the completed work, then make it observable

Adopter decision	Evidence that tests the decision	How to design the criterion
Completed deliverable	What must be returned, saved or changed?	Write a contract from real valid cases
Unacceptable errors	Wrong targets, false completion, unauthorized change	Set withhold / reject rules by risk
Success and resource limits	Verifier, time, cost and retry records	Test boundary cases; freeze thresholds

A design guide, not a recommendation to adopt this study’s thresholds as a business SLO.

The adopter defines completion; verifiers check it


Adopter design guide / this diagnostic contract

What observable evidence counts as completion?

Our task: return final JSON for every read-only inbox item

A fixed synthetic sandbox contract, not customer delivery or an order mutation

Status request	Look up the target; consume its status
Unclear target	needs_clarification
Cancel / refund	unsupported · no mutation tool
Final shape	One items entry per input ID



Canonical	Valid trial + post-run pass + runtime final acceptance
Strict	Canonical + designated helper path / consumption + zero wrong or extra lookups / false completion

Initial helper accuracy is separate from completion: the root may recheck the request and correct the helper.

The adopter defines completion; verifiers check it

Adopter design guide / this diagnostic contract


What observable evidence counts as completion?

Two ordered responses, one pool-level score

One pool from two retained Jev responses; an observed example of the scoring method

Typed / numeric validator: response validity, not answer correctness.

Candidate-aligned means, two orders



Highest mean score

A
argmax = A
2.325

Pre-run oracle grades

A: 3 B: 0 C: 2 D: 0

Selected A: grade 3

Matches the highest grade

1 / 1 Pool correctness

Two orders measure one pool, not two independent samples; ties use the frozen hash rule.

Resample related items together by source; the 95% interval shows uncertainty in the accuracy difference.

Score output errors / fallback count as wrong; infrastructure invalidity or non-execution is not zero performance.

One pointwise example. The listwise reference averages correctness across two orders: 0 / 0.5 / 1.

Different data measure different questions

Final native evidence · editorial interpretation

Items, calls or independent tasks: what is counted?

Different datasets do not share one success denominator

Source / author	Visibility / split	Denominator / independent unit	Gold authority
Authored judgment panel 45 source clusters	Selection 18 · Test 27 Public inputs; gold separate	160 / 240 states 216 primary test states	Rule oracle Read after test unseal
I · intent and target Authored by the designer	Unsealed diagnostic data Admission 4 · Main 40 families	20 / 220 items E2E: 36 items ≠ 3 tasks	Authored labels Not held out
E · inbox tasks Separate 20-task set	12 natural + 6 injected + 2 reserve Fresh containers per arm	U7: 12 tasks × 3 × 2 U8c: 6 tasks × 3 faults × 2	Independent verifier Separate from judge
External X1 / X2 Human labels / Mind2Web	Text conversion / K4 selection X2: 80 per official split	239 states / 240 pools 133 / 94 clusters	Source gold retained Score after all attempts

Different denominators for states, items, pools, tasks and calls. Compare success within the same unit.

Different data measure different questions

Final native evidence · editorial interpretation

Items, calls or independent tasks: what is counted?

How are related states grouped by source?

Verdict panel · 45 source clusters

Score · 80 controlled pools

Selection 18 / 160 states

144 primary + 16 out-of-envelope

Selection 32 + Test 48

Not all 80 pools are held out

Test 27 / 240 states

216 primary + 24 out-of-envelope

12 natural pools · best-of-4

Fixed after generation; no regeneration to force success

40:60 is a source-cluster ratio, not proof of random assignment.

Public sibling states can reveal the m condition. This is not fully blind evaluation.

A 2×2 design separates contradiction from missing evidence

Verdict panel / controlled changes to structured states

Can contradiction and missing evidence be tested separately?

Vary two factors within one source: claim, observation and contract

		c · request / contract / evidence conflict	
		c = 0 · absent	c = 1 · present
m · required evidence			
m = 0 Present		Neither defect c0m0	Contradiction only c1m0
m = 1 Missing		Missing evidence only c0m1	Both defects c1m1

Gold authority

Rule oracle

Check the designed factors and rules

Model responses are evaluated

They do not generate the gold

Not a human audit of real-world truth

90 × 4 = 360

90 quadruplets · primary states

+ 40

Out-of-envelope states

= 400


All selection + test states

45 source clusters; four related states are not four independent samples.

Who reads which evidence before the next step?

Final native evidence · editorial interpretation

Who produces the evidence, and who verifies it?



Playback and film do not own scores. E admission audit / delegated-playback gates follow.


Different prerequisites connect the 27 registered units

Final native evidence · editorial interpretation

Does passing admission establish main-study performance?

P · verdicts / calibration / consistency / latency

9 registered units · arrows encode prerequisites, not concurrent execution.



Attempt closure and analysis completion are different states.
Only after U2 → U2s → U3 close does explicit unseal approval permit fixed analysis.


Different prerequisites connect the 27 registered units

Final native evidence · editorial interpretation

Does passing admission establish main-study performance?

G/P · selection / P · intent / X · external transfer

9 registered units · X uses the P driver; it is not a separate concurrent runner.



Shared X gates: core P closure / unseal / run order. Read external labels only after its attempts close.
Preserve non-discriminative U5 pools; do not regenerate candidates until selection succeeds.


Different prerequisites connect the 27 registered units

Final native evidence · editorial interpretation

Does passing admission establish main-study performance?

E · tasks, recovery, intent handoff, C repetitions

9 registered units · same-source admission, audit and delegated playback gate downstream runs.




Before C admission: frozen τ / U2c gate / explicit C and budget conditions. The 27 IDs are not model wins.

Freeze policy on Selection; open gold after execution

Final native evidence · editorial interpretation

When was the policy frozen, and when was gold opened?




Choice T: Astra 8 / Jev 1.4 · $\tau=0.50$. T is post-hoc calibration, not generation temperature.

AUROC / C use raw $q=\max P$. Keep q distinct from separately reported native confidence.

Freeze policy on Selection; open gold after execution

Final native evidence · editorial interpretation

When was the policy frozen, and when was gold opened?



Public sibling states could reconstruct m (216/216); models received one state at a time.
There is no Test → T / τ refit path. Hashes and access rules are not encryption or OS isolation.

Before the numbers: what each metric answers

Definitions / not new experimental results

Start with units, direction and criteria

Comparison charts / difference and uncertainty

Δ · pp

How far apart are the two accuracy rates?

Δ = Jev minus Astra, in percentage points, not relative percent change.

95% CI

How uncertain is the estimated difference?

An interval from resampling source groups, not a probability for one response.

NI · margin

Is the loss within the predeclared margin?

Here: the CI lower bound must exceed -5 or -10 pp. Not a business SLO.

Source cluster

What counts as an independent source unit?

Related states from one source stay together; call counts are not independent n.

Before the numbers: what each metric answers

Definitions / not new experimental results

Start with units, direction and criteria

Probability signals / calibration, error ranking and acceptance

ECE ↓

Lower is better, not a trend

How well do stated probabilities match observed accuracy?

Valid-output bins: sample-weighted absolute gaps between mean q and accuracy.

AUROC ↑

Higher is better, not a trend

Does the signal rank incorrect responses first?

Here, $1-q$ ranks errors ahead of correct answers; 0.5 is the random-ranking reference.

Coverage / Risk

How many are accepted, and how many accepted answers are wrong?

Coverage = accepted / planned; risk = incorrect accepted / all accepted.

AURC ↓

Lower is better, not a trend

How much error accumulates as acceptance expands?

Area under the risk–coverage curve; lower error burden, not a safety guarantee.

Before the numbers: what each metric answers

Definitions / not new experimental results

Start with units, direction and criteria

Repeated execution / success, failure and measurability

pass@1 / @2 / ^2

Per attempt, at least once, or both times?

Here: mean per-attempt success / either fresh repeat / both; not best-of-four selection.

canonical / strict

Beyond completion, were the required path and lookup rules met?

Canonical: valid run, task checks, final acceptance. Strict adds path / lookup / delivery rules.

Failure / Invalid

A model failure, or evidence that cannot support measurement?

Valid errors are outcomes; infra-invalid / unrun is not 0%. Output errors follow metric rules.

Admission / Test

A preflight check of the execution path, or a main-study outcome?

Admission gates the next run; its passes do not enter the main-study success rate.

Close mean accuracy, with a wide interval

Measured · native results retained

Is the accuracy loss within the predeclared margin?

Choice · primary analysis, 216 states

A · Astra xhigh

157/216

B · Jev 1.13.0

155/216

Accuracy difference $\Delta = \text{Jev} - \text{Astra}$ · percentage points (pp)

Negative: lower Jev accuracy

Positive: higher Jev accuracy



Dot: estimated difference; 95% CI: uncertainty

Solid 0: no difference

Dashed: NI -5 pp

$\Delta -0.93$ pp · 95% CI [-9.72, +7.41]

Mean difference -0.93 pp; the wide CI does not establish -5 pp non-inferiority

4×4 order rotation; one invalid model output counts as wrong; 27-source bootstrap, 2,000 replicates

A large loss on the joint judgment

Panel measured · Prior source5 Harbor follow-up: closed, not run

Were both conditions correct together?

Noul · share correct on both contradiction and missing evidence

A · Astra xhigh

145/216


B · Jev 1.13.0

55/216

Accuracy difference $\Delta = \text{Jev} - \text{Astra}$ · percentage points (pp)

Negative: lower Jev accuracy

Positive: higher Jev accuracy



Dot: estimated difference; 95% CI: uncertainty

Solid 0: no difference

Dashed: NI -5 pp

$\Delta -41.67$ pp · 95% CI [-48.15, -35.19]

The whole CI lies below -5 pp; individual accuracy cannot replace joint accuracy

216 states · 2,000 cluster-bootstrap replicates · 27 source clusters

A small loss on the fixed intent panel

Intent / target panel • 220 items / 40 families

How many items have both the right intent and target?

Joint intent and target • 220 items / 40 families

A • Astra xhigh

220/220

B • Jev 1.13.0

216/220

Accuracy difference $\Delta = \text{Jev} - \text{Astra}$ • percentage points (pp)

Negative: lower Jev accuracy

Positive: higher Jev accuracy



Dot: estimated difference; 95% CI: uncertainty

Solid 0: no difference

Dashed: NI -5 pp

$\Delta -1.82$ pp • 95% bootstrap CI $[-3.65, -0.45]$

CI lower bound exceeds -5 pp: non-inferior on this panel, not evidence of generalization

Designer-authored panel; 40-family percentile bootstrap, 2,000 replicates

Selection accuracy is not pool discrimination

Post-hoc corrected analysis of retained responses

Was a candidate with the highest oracle grade selected?

Score • controlled-pool top-1 accuracy

A • Astra xhigh

79/80


B • Jev 1.13.0

59/80

Accuracy difference $\Delta = \text{Jev} - \text{Astra}$ • percentage points (pp)

Negative: lower Jev accuracy

Positive: higher Jev accuracy



Dot: estimated difference; 95% CI: uncertainty

Solid 0: no difference

Dashed: NI -10 pp

$\Delta -25.00$ pp • 95% CI [-35.23, -16.25]

The whole CI lies below -10 pp; Jev selects less accurately in these pools

Retained-response parser-corrected reanalysis; fixed denominator / grades; 45 clusters

Selection accuracy is not pool discrimination

Post-hoc corrected analysis of retained responses

Was a candidate with the highest oracle grade selected?

12 natural pools · best-of-4	Correct selection	gap_closed
Astra pointwise	12/12	Not measurable
Jev pointwise	12/12	Not measurable
First / random / oracle	All 12/12	Non-discriminative

gap_closed: how much of the random-to-oracle gap was closed?

Denominator: oracle coverage – random = 0; accuracy Δ = 0 pp; neither equivalence nor superiority

Fixed candidates generated once by Astra; no regeneration or selector superiority claim

More correct labels, but many human successes rejected

Measured · native results retained

Are failures caught without rejecting genuine successes?

CUAVerifierBench · external authors and human labels

Test verdict transfer beyond the authored panel

260 states – 21 human ties = 239 states

133 source clusters; not the original benchmark score

Text projection: actions, thoughts and notes

Screenshots and actual tool-result contents are absent

Policy remains fixed from internal selection

No retuning T or τ on external test data

Preserve trajectory / human-label limits; not a live computer-use success rate

More correct labels, but many human successes rejected

Measured · native results retained

Are failures caught without rejecting genuine successes?

Binary verdicts on external text-only states

A · Astra xhigh

125/239

B · Jev 1.13.0

136/239

Accuracy difference $\Delta = \text{Jev} - \text{Astra}$ · percentage points (pp)

Negative: lower Jev accuracy

Positive: higher Jev accuracy



Dot: estimated difference; 95% CI: uncertainty

$\Delta +4.60$ pp · 95% CI [+0.41, +9.36]

Positive relative difference; human-success rejections: Astra 113 / Jev 97 of 119

260 minus 21 ties = 239 states · no screenshots or tool-result contents · 133 source clusters

Lower selection accuracy on the same four candidates

Post-hoc corrected analysis of retained responses

One-step selection or completion of a whole web task?


Three test splits

cross-task / website / domain

Controlled pool K = 4

One gold + three negatives · two orders

Score selection



240 steps · 80 per split

Candidate selection for a step, not full web-navigation success

Pool construction controls difficulty; not the original full-ranking benchmark

Show transforms, IDs and hashes only; no raw questions or HTML

No threshold or rubric tuning on external data

Lower selection accuracy on the same four candidates

Post-hoc corrected analysis of retained responses

One-step selection or completion of a whole web task?

Mind2Web • the same four candidates, 240 steps

A • Astra xhigh

226/240

B • Jev 1.13.0

196/240

Accuracy difference $\Delta = \text{Jev} - \text{Astra}$ • percentage points (pp)

Negative: lower Jev accuracy

Positive: higher Jev accuracy



Dot: estimated difference; 95% CI: uncertainty

Solid 0: no difference

Dashed: NI -10 pp

$\Delta -12.50$ pp • 95% CI $[-17.22, -8.04]$

CI below zero, crossing -10 pp; operational justification for the registered margin is weak

Retained-response parser-corrected reanalysis; fixed denominator / grades; 94 clusters


Shorter response waits for the same judgments

Measured · native results retained

How long did matched judgment calls take?

Observed response times across 120 matched pairs

Vertical axis: share of calls completed within that time



Astra · xhigh

10.446 s

Median · 120 calls

Jev · 1.13.0

0.253 s

Median · 120 calls

Judge-call waiting time, separate from whole-task runtime.

Shorter response waits for the same judgments

Measured · native results retained

How long did matched judgment calls take?

Concurrent dispatch · 120/120 comparable pairs

Concurrent pairs

120

Source clusters

27

Median of paired B–A verdict times · seconds



Dot: estimated difference; 95% CI: uncertainty

Solid 0: no difference

No NI margin

Δ -10.20 s · 95% CI [-11.34, -8.81]

Uncertainty of the verdict-time difference; not per-call spread or whole-task time

median($B_i - A_i$), not median(B)–median(A) · 27 source clusters

Separate call latency from cost accounting

Ledger after E7 U7r1 · 2026-09-28 00:34 KST

Are estimated costs separated from actual invoices?

Recorded calls, tariff estimates and invoices are distinct records

Recorded calls

Jev usage

Astra 3,753 · Jev 2,267

5,324,918 input tokens

Includes failed, stopped and recovery calls

known 2,263 · unknown 4 calls

\$0.223646556

\$0.004200000

unknown

Input-tariff estimate

Unknown-usage reserve

Actual invoice

Committed \$0.227846556; operating cap \$2.35. Neither is an invoice.

Do not convert subscription Astra into PAYG savings. Full physical-call coverage is unknown.


Noul / Jev: 23/24 matched answers, 6/24 correct twice

24-state repetition diagnostic; native results retained

Repeated answers, or repeated correct answers?

24 states per primitive; two fresh calls on the same input

Beyond cost and latency: are repeated answers also correct?



Noul / Jev: 23/24 consistent, only 6/24 correct twice. Repeated errors also count as consistent.


Noul / Jev: 23/24 matched answers, 6/24 correct twice

24-state repetition diagnostic; native results retained

Repeated answers, or repeated correct answers?

Meaning-preserving variants: states where Jev changed its answer

Order reversal and paraphrase compared with the original, separate from fresh repeats



Low flip means less response sensitivity here, not high accuracy or E2E pass².


Next: separately test whether reported probabilities distinguish errors

How well do verdict probabilities identify errors?

Measured · native results retained

Are calibration and error ranking the same capability?

Coverage and error risk ranked by raw $q = \max P$



Descending q ; coverage: accepted / all 216 states; risk: incorrect / accepted

How well do verdict probabilities identify errors?

Measured • native results retained

Are calibration and error ranking the same capability?

Different questions for the same probability signal

	Astra • xhigh	Jev • 1.13.0
Raw ECE ↓ Do stated probabilities match observed accuracy?	0.260	0.075
Calibrated ECE ↓ How much mismatch remains after calibration?	0.109	0.077
AUROC ↑ Does the signal rank incorrect responses first?	0.879	0.785

On this panel: lower Jev ECE, higher Astra AUROC; no single combined ranking


ECE: ten equal-width bins; AUROC detects errors with 1-q; q=max P ≠ native Confidence

Valid probabilities: Astra 216 / Jev 215; T: 8.0 / 1.4, frozen on selection

Separate conditional acceptance from actual fallback

$\tau=0.50$ frozen on selection

Which responses are accepted, and which go to review?



Choose τ on selection → evaluate the same policy on test

Do not equate Jev confidence with q

Accept Jev only when validity and q pass; otherwise obtain a separate Astra verdict.

Separate conditional acceptance from actual fallback

$\tau=0.50$ frozen on selection

Which responses are accepted, and which go to review?

Conditional acceptance on the panel

Astra · xhigh

157/216

Jev + Astra fallback

160/216

Jev coverage 190/216 · $\tau = 0.5$

$\Delta +1.39$ pp · 95% CI $[-6.48, +9.26]$ pp

Conditional policy computed on the same test responses

Response selection on the panel, not the count of fallbacks in a live run


The CI includes zero: superiority unproven. Live fallback utility needs a separate test.

Distinguish paired execution from admission to the next slot

Final native evidence · editorial interpretation

Did the arms start under the same comparison rules?

A global lock between units; isolated arms within each slot



Dispatch >1s: infra-invalid; agent-start skew >30s: exclude paired latency

Separate sync and semantic failures. The two U7 r1 pairs stay in strict/reliability counts.

All three policies completed the tasks under the final conditions

Final native evidence · editorial interpretation

Success at least once, or in both fresh repetitions?

12 independent tasks · three conditions · two fresh repetitions

One filled dot = one strict-success task; columns are repetitions, rows are verifier policies

	Repeat 1													Repeat 2												
A · Astra	●	●	●	●	●	●	●	●	●	●	●	12/12		●	●	●	●	●	●	●	●	●	●	12/12		
B · Jev	●	●	●	●	●	●	●	●	●	●	●	12/12		●	●	●	●	●	●	●	●	●	●	12/12		
C · conditional	●	●	●	●	●	●	●	●	●	●	●	12/12		●	●	●	●	●	●	●	●	●	●	12/12		

72/72 complete, but not 72 independent tasks

Every arm: pass@1 = pass@2 = pass^2 = 12/12; per-attempt mean / at least once / both

C: 24 accepts, zero actual fallbacks. Execution worked; saturated success did not distinguish policy utility.

Both verifiers caught injected errors; Astra repaired them

Final native evidence · editorial interpretation

Did detecting the fault lead to repair and completion?

Injected condition	Required repair	Completed trials
Contradiction only	Fix answer conflicting with evidence	6 tasks × A/B
Missing evidence only	Supply and use required evidence	6 tasks × A/B
Contradiction + missing	Resolve both in one repair round	6 tasks × A/B

36 Choice trials · paired A/B comparison within each defect combination

Both verifiers caught injected errors; Astra repaired them

Final native evidence · editorial interpretation

Did detecting the fault lead to repair and completion?

Observed Choice recovery / 18 trials per arm

	A · Astra xhigh	B · Jev 1.13.0
Initial negative verdict	18/18	18/18
Recovered / strict success	18/18	18/18
Astra replans per trial	1	1

All 36: negative feedback consumed, Astra repair, then strict completion; false completion 0

18 pairs / six sources; recovery difference 0; no CI; descriptive comparison


Paired B–A runtime median –18.17 s; every trial below 180 s

A planned injection experiment, not general recovery reliability or a causal resource-budget effect.

Resolve a 12-request order inbox · Choice

Completion verification: injected error, repair and final return


A · Astra xhigh



Clip 102.50 s; native runtime 101.60 s

ATIF-derived · not raw PTY or simultaneous execution · film 4× / source clocks · scores and latency from native results

B · Jev verdict · Astra root



Clip 94.93 s; native runtime 94.03 s

An unsupported refund claim and missing lookup evidence

Source7 Noul execution / post-hoc analysis

What evidence was available to the verifier?

12 synthetic requests; Astra xhigh executes and repairs in both arms

Four inputs sent to the verifier

task_contract

Read-only: look up, clarify or mark unsupported

original_request

12 requests: status, refund, correction, multiple orders

candidate_output

Candidate JSON; the first candidate has an injected fault

tool_observations

Actual lookup inputs and results; initially absent

c: contradiction with the request, contract or observations; m: missing evidence

The task answer key and hidden reasoning are separate from verifier input

Initial injected fault

Request	Refund / T-846
Allowed tools	Read-only lookup No mutation tool
Injected	status: refunded disposition: answered Required lookups absent

First verdict	Both: c=1, m=1
Astra consumed the negative feedback	

An unsupported refund claim and missing lookup evidence

Source7 Noul execution / post-hoc analysis

What evidence was available to the verifier?

Five actual lookups and Astra’s contract-compliant candidate

Five status items in one lookup batch, not five tool calls

T-503	processing	Refund request	unsupported status=null
T-218	shipped		
T-691	on_hold	5 status requests	Match actual lookups
T-135	delivered	Other 7 requests	Unsupported or needs_clarification
T-284	out_for_delivery		

12/12

Correct items in both final candidates, separate from delivery success

Matched item meanings, contract and lookup states; JSON formatting and call IDs are not byte-identical.

Noul: why a correct candidate was withheld


Source7 Noul execution / post-hoc analysis

Why was the correct candidate not returned?

Both candidates: 12/12 correct; all five required lookups were in the verifier input

Astra executes and repairs in both arms; A uses Astra verification, B uses Jev

Missing-evidence probability (m)



Contradiction (c): A 0.00 / B 0.38; both below 0.50

A: c=false / m=false; supported; JSON returned

The delivery gate in arm B

Jev output	$m = 0.63 \geq 0.50$ c=false / m=true
Adapter verdict	insufficient_evidence passed=false
GEODE policy	Continuation budget: 2/2 used Candidate retained in pending_text Final text=""; delivery withheld

Despite complete evidence: m=0.60, then 0.63; false rejection under this task contract
No item-level rationale returned; Reflection is GEODE-authored. The triggering item or representation is unidentified.

Noul: why a correct candidate was withheld

Source7 Noul execution / post-hoc analysis

Why was the correct candidate not returned?

Separate candidate correctness from strict completion: four sources, 11 complete pairs

Final candidates correct: 11/11 in both arms; strict/recovered: A 11/11, B 0/11; B withheld 11/11

		A · Astra			B · Jev		
		c only	m only	Both	c only	m only	Both
telescopes	3 pairs	●	●	●	□	□	□
badminton	3 pairs	●	●	●	□	□	□
waterpurifiers	3 pairs	●	●	●	□	□	□
strollers	2 pairs	●	●	-	□	□	-

Filled: strict complete; open: withheld; -: incomplete pair excluded


Planned 36; observed 24 (23 valid / 1 invalid); 12 unrun; original run invalidated, primary null

11 pairs are not 11 independent sources; no CI / NI / population claim; correct candidate ≠ recovered

Resolve a 12-request order inbox · Noul


Five required lookups completed; Jev’s missing-evidence verdict withheld delivery

A · Astra xhigh



Source-media time; not task latency

B · Jev verdict · Astra root



Source-media time; not task latency

ATIF-derived · not raw PTY or simultaneous execution · film 4× / source clocks · scores and latency from native results

Final inbox execution completed both judgments and delivery


Before / after contract correction remain separate evidence lineages

Are helper correctness and task completion counted separately?

Final conditions / source7: separate item judgments from task completion


Helper · intent and target separately

36 items per arm · initial = final



Independent tasks · strict completion

Three trials per arm · one repeat



X: correct items (0–36)

One filled dot = one strict-success trial

Three observed trials per arm: no rejudgment, correction, wrong-target lookup or false completion


Δ strict = 0 pp; three pairs; no CI; descriptive, no superiority / equivalence test

Three author-visible synthetic inboxes; one repeat; final JSON, not customer delivery

Classify 12 Korean requests · Choice

Intent and target helper; Astra remains the final verifier


A · Astra xhigh



Clip 93.23 s; native runtime 92.55 s

ATIF-derived · not raw PTY or simultaneous execution · film 4× / source clocks · scores and latency from native results

B · Jev helper · Astra root



Clip 77.07 s; native runtime 76.25 s

The offloading scope supported by these results

Final native evidence · editorial interpretation

How far does the evidence support delegation?

Translate observations into decisions within the measured scope

Decision	Observation and implication	Decision and adoption conditions
Fixed intent	Astra 220/220 · Jev 216/220; E2E 3/3 each	Validate on new inboxes and sources
Choice completion	Small mean difference, wide interval	An operational loss tolerance
Noul / delivery	Astra 145/216 · Jev 55/216 Post-hoc: Jev correct candidates 11/11; delivery 0/11	No support for standalone replacement
Conditional C	E7 U7, two repeats: C 24 accept / 0 fallback	Utility when errors actually escalate


This is an evidence-bounded interpretation, not a product-wide performance or reliability guarantee.

Read the study, then follow the evidence

Pinned public evidence / artifact 3bcf4044

Original results, corrections and post-hoc diagnoses kept separate

Public evidence separates the pilot, original results and post-hoc analyses



Earlier pilot Implementation / judgment-latency diagnostic; separate denominators

Main study v3 27-unit index; pinned specs, results, denominators and analyses

Additive evidence: parser correction and 11-pair Noul diagnosis

Existing pilot · study/publication contract

Main study / pinned public README

github.com/mangowhoisccloud/geode

github.com/mangowhoisccloud/geode-eval-artifacts

Published commit 3bcf4044; original and corrected analyses have separate paths; withheld inputs disclosed

Checking a result is not a new experiment

Pinned public evidence / artifact 3bcf4044

Recomputing retained evidence, or making new model calls?

Shared entry: read README, commit, manifest and SHA; check integrity


Reader path	Inputs and interpretation	Experiment calls
Recompute fixed results	Predictions / gold / clusters / pinned analyzer Recompute: rerun analysis on fixed inputs, with no new model calls	None
New run on own account	Own account / budget / new freeze / run ID A separate, condition-bound run; numbers may differ	Required
Trajectory / Replay	Public trajectories / selected Replay / source clock and speed Replay: view a recorded run; check omissions. Native results own scores and latency.	None
X2 · Mind2Web	Official source → pinned converter → hashes; no republished test text/pools If recomputation inputs are missing, disclose the unsupported scope. Public recipes: U2, U4, corrected U4/X2 and Noul post-hoc analysis; not private-environment validation.	

Give an agent the same evidence trail

Pinned public evidence / artifact 3bcf4044

Original results, corrections and post-hoc diagnoses kept separate

A read-only review follows one result all the way to its evidence



Read-only review: no experiment APIs, auth changes, paid runs or T/τ refitting

Report commit and input hashes, metric / denominator matches, and unverified inputs or execution scope

Tracing the interpretation to sources and execution records

Question	Evidence owner	Inspect
How much of Jev is public?	API contract · vendor account Primitives · System One · RLCD	API · SDK Launch · Interview
Where did time and errors change?	Our separate measurement cohorts Natural runs: 6 pairs · Fixed inputs: 3 states · Fault injection: 2 conditions	Runs / metrics
Did the label trigger a repair?	Pinned code and execution trace Verdict → feedback → Astra → oracle	Paired Replay Code / records
Background: RSI taxonomy · Qwen/CLM · dual-process theory · Reflexion (Shinn et al., NeurIPS 2023) → public study evidence		

Production references

REFERENCES → APPLICATION IN THIS FILM

Beyond Pass Rates · GEODE ↗

Restrained typography; claims followed by execution evidence.

Stanford CS231n · Lecture 2, pp. 7–8 ↗

Expand one example from input to representation and candidate labels.

Anthropic · Tracing the thoughts... ↗

Connect intermediate states to outcomes in one visual.


What's next after RLHF? · AI Engineer ↗

Diogo Almeida: questions and counterexamples clarify training objectives.

Anthropic · LangGraph · OpenAI ↗

Workflow gates, feedback edges and repository-enforced invariants.

ASSETS USED ON SCREEN



TypeSafe symbol · GEODE Geodi

가 Aa

[IBM Plex Sans / KR ↗](#)

Titles, body and diagrams · code: Plex Mono

[Design references and font provenance ↗](#)

Faster calls are supported. Blanket delegation is not.

Final native evidence · editorial interpretation

Original runs preserved; corrections labeled separately

Bound delegation by the measured gains and losses

Faster verifier calls	−10.20 s Not whole-task speedup 120 pairs · paired median · CI [−11.34, −8.81] s
Structured intent is promising	216/220 Final inbox E2E: 3/3 per arm Jev; Astra 220/220; −5 pp NI met on this panel
Controlled Score selection gap	−25.00 pp Insufficient support for standalone replacement Noul joint −41.67 pp; external selection −12.50 pp
Jev strict completion after repair	0/11 False rejections also need verification Candidates correct 11/11; withheld 11/11; post-hoc, four sources / 11 pairs

Before adoption: validate helpers on new inputs; measure final-verifier false accepts and false rejects separately.